

## 6.5.3 Choosing Among Models Using the Validation Set Approach and Cross-Validation

May 18, 2018

```
In [1]: # conventional way to import pandas
import pandas as pd
# conventional way to import seaborn
import seaborn as sns
# conventional way to import numpy
import numpy as np

from sklearn import metrics
import matplotlib.pyplot as plt
```

```
data = pd.read_csv("https://vincentarelbundock.github.io/Rdatasets/csv/ISLR/Hitters.csv")
```

```
data.head()
```

```
Out[1]:
```

	AtBat	Hits	HmRun	Runs	RBI	Walks	Years	CAtBat	CHits	\
-Andy Allanson	293	66	1	30	29	14	1	293	66	
-Alan Ashby	315	81	7	24	38	39	14	3449	835	
-Alvin Davis	479	130	18	66	72	76	3	1624	457	
-Andre Dawson	496	141	20	65	78	37	11	5628	1575	
-Andres Galarraga	321	87	10	39	42	30	2	396	101	

	CHmRun	CRuns	CRBI	CWalks	League	Division	PutOuts	\
-Andy Allanson	1	30	29	14	A	E	446	
-Alan Ashby	69	321	414	375	N	W	632	
-Alvin Davis	63	224	266	263	A	W	880	
-Andre Dawson	225	828	838	354	N	E	200	
-Andres Galarraga	12	48	46	33	N	E	805	

	Assists	Errors	Salary	NewLeague
-Andy Allanson	33	20	NaN	A
-Alan Ashby	43	10	475.0	N
-Alvin Davis	82	14	480.0	A
-Andre Dawson	11	3	500.0	N
-Andres Galarraga	40	4	91.5	N

After listing the data we can see that some have missing data for their Salary. Next drop all the rows that contain NaN data.

```
In [2]: data = data.dropna()
data.index.name = 'Player'
data.head()
```

```
Out[2]:
```

	AtBat	Hits	HmRun	Runs	RBI	Walks	Years	CAtBat	CHits	\
Player										
-Alan Ashby	315	81	7	24	38	39	14	3449	835	
-Alvin Davis	479	130	18	66	72	76	3	1624	457	
-Andre Dawson	496	141	20	65	78	37	11	5628	1575	
-Andres Galarraga	321	87	10	39	42	30	2	396	101	
-Alfredo Griffin	594	169	4	74	51	35	11	4408	1133	

	CHmRun	CRuns	CRBI	CWalks	League	Division	PutOuts	\
Player								
-Alan Ashby	69	321	414	375	N	W	632	
-Alvin Davis	63	224	266	263	A	W	880	
-Andre Dawson	225	828	838	354	N	E	200	
-Andres Galarraga	12	48	46	33	N	E	805	
-Alfredo Griffin	19	501	336	194	A	W	282	

	Assists	Errors	Salary	NewLeague
Player				
-Alan Ashby	43	10	475.0	N
-Alvin Davis	82	14	480.0	A
-Andre Dawson	11	3	500.0	N
-Andres Galarraga	40	4	91.5	N
-Alfredo Griffin	421	25	750.0	A

Change the string values into dummy values

```
In [3]: dummieVariables = pd.get_dummies(data[['League', 'Division', 'NewLeague']])
dummieVariables.info()
print(dummieVariables.head())

y = data.Salary
X_ = data.drop(['Salary', 'League', 'Division', 'NewLeague'], axis=1).astype('float64')

X = pd.concat([X_, dummieVariables[['League_N', 'Division_W', 'NewLeague_N']]], axis=1)
X.head()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 263 entries, -Alan Ashby to -Willie Wilson
Data columns (total 6 columns):
League_A      263 non-null uint8
League_N      263 non-null uint8
Division_E     263 non-null uint8
Division_W     263 non-null uint8
NewLeague_A    263 non-null uint8
NewLeague_N    263 non-null uint8
```

dtypes: uint8(6)

memory usage: 3.6+ KB

	League_A	League_N	Division_E	Division_W	NewLeague_A	\
Player						
-Alan Ashby	0	1	0	1	0	
-Alvin Davis	1	0	0	1	1	
-Andre Dawson	0	1	1	0	0	
-Andres Galarraga	0	1	1	0	0	
-Alfredo Griffin	1	0	0	1	1	

	NewLeague_N
Player	
-Alan Ashby	1
-Alvin Davis	0
-Andre Dawson	1
-Andres Galarraga	1
-Alfredo Griffin	0

Out[3]:

	AtBat	Hits	HmRun	Runs	RBI	Walks	Years	CAtBat	\
Player									
-Alan Ashby	315.0	81.0	7.0	24.0	38.0	39.0	14.0	3449.0	
-Alvin Davis	479.0	130.0	18.0	66.0	72.0	76.0	3.0	1624.0	
-Andre Dawson	496.0	141.0	20.0	65.0	78.0	37.0	11.0	5628.0	
-Andres Galarraga	321.0	87.0	10.0	39.0	42.0	30.0	2.0	396.0	
-Alfredo Griffin	594.0	169.0	4.0	74.0	51.0	35.0	11.0	4408.0	

	CHits	CHmRun	CRuns	CRBI	CWalks	PutOuts	Assists	\
Player								
-Alan Ashby	835.0	69.0	321.0	414.0	375.0	632.0	43.0	
-Alvin Davis	457.0	63.0	224.0	266.0	263.0	880.0	82.0	
-Andre Dawson	1575.0	225.0	828.0	838.0	354.0	200.0	11.0	
-Andres Galarraga	101.0	12.0	48.0	46.0	33.0	805.0	40.0	
-Alfredo Griffin	1133.0	19.0	501.0	336.0	194.0	282.0	421.0	

	Errors	League_N	Division_W	NewLeague_N
Player				
-Alan Ashby	10.0	1	1	1
-Alvin Davis	14.0	0	1	0
-Andre Dawson	3.0	1	0	1
-Andres Galarraga	4.0	1	0	1
-Alfredo Griffin	25.0	0	1	0

Now we will split the data into a training and testing set

```
In [4]: from sklearn.model_selection import train_test_split
        X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=1)
```

bla bla

```

In [5]: import itertools
import statsmodels.api as sm
max_feature=18

# Functions found at "https://github.com/qx0731/ISL_python/blob/master/Chapter_6_sec_6
def validateRSS(y_train, X_train, y_test, X_test, predictors_list):
    model = sm.OLS(y_train, X_train[list(predictors_list)]).fit()
    RSS = ((model.predict(X_test[list(predictors_list)]) - y_test) ** 2).sum()
    return {'Model':model, "RSS":RSS}

def validateBestModel(y_train, X_train, y_test, X_test, K):
    results = []
    for c in itertools.combinations(X_train.columns, K):
        results.append(validateRSS(y_train, X_train, y_test, X_test, c))
    allModels = pd.DataFrame(results)

    bestModel = allModels.loc[allModels["RSS"].argmin()]
    return bestModel

def validateForwardStepwiseSelection(y_train, X_train, y_test, X_test, predictors_list):
    remaining_predictors = [p for p in X_train.columns if p not in predictors_list]
    results = []
    for p in remaining_predictors:
        results.append(validateRSS(y_train, X_train, y_test, X_test, predictors_list+[p]))

    allModels = pd.DataFrame(results)
    bestModel = allModels.loc[allModels['RSS'].argmin()]
    return bestModel

def validateBackwardStepwiseSelection(y_train, X_train, y_test, X_test, predictors_list):
    results = []
    for combo in itertools.combinations(predictors_list, len(predictors_list)-1):
        results.append(validateRSS(y_train, X_train, y_test, X_test, combo))

    allModels = pd.DataFrame(results)
    bestModel = allModels.loc[allModels['RSS'].argmin()]
    return bestModel

C:\Users\au479931\AppData\Local\Continuum\anaconda3\lib\site-packages\statsmodels\compat\pandas
from pandas.core import datetools

bla bla

In [6]: validationModels = pd.DataFrame(columns=["RSS", "Model"])
for i in range(1, (max_feature+1)):

```

```
validationModels.loc[i] = validateBestModel(y_train, X_train, y_test, X_test, i)
```


```
validationModelsForward = pd.DataFrame(columns=["RSS", "Model"])  
predictors_list = []  
for i in range(1, len(X.columns)+1):  
    validationModelsForward.loc[i] = validateForwardStepwiseSelection(y_train, X_train,  
    predictors_list = validationModelsForward.loc[i]["Model"].model.exog_names
```

C:\Users\au479931\AppData\Local\Continuum\anaconda3\lib\site-packages\ipykernel\_launcher.py:18

C:\Users\au479931\AppData\Local\Continuum\anaconda3\lib\site-packages\ipykernel\_launcher.py:29

bla bla


```
In [7]: plt.figure()  
plt.plot(validationModels["RSS"])  
plt.xlabel('Predictors')  
plt.ylabel('RSS')  
plt.show()
```




bla bla

```
In [8]: plt.figure()  
plt.plot(validationModelsForward["RSS"])
```

```
plt.xlabel('Predictors')
plt.ylabel('RSS')
plt.show()
```



```
In [16]: plt.figure()
plt.plot(validationModels["RSS"], 'b')
plt.plot(validationModelsForward["RSS"], 'y')
plt.xlabel('Predictors')
plt.ylabel('RSS')
plt.show()
```



Now we've tried to validate best subset and forward stepwise selection, let's continue with Cross-Validation

```
In [9]: k = 10
        np.random.seed(seed = 21)
        train_index = np.random.choice(k, size = len(y), replace = True)
        crossValidationErrors = pd.DataFrame(columns=range(1,k+1), index=range(1,len(X.columns)))

        bla bla
```

```
In [10]: crossValidationModels = pd.DataFrame(columns=["RSS", "Model"])
         for j in range(1,k+1):
             predictors_list = []
             for i in range(1,len(X.columns)+1):
                 crossValidationModels.loc[i] = validateForwardStepwiseSelection(y[train_index], X[predictors_list])


                 crossValidationErrors[j][i] = crossValidationModels.loc[i]["RSS"]
                 predictors_list = crossValidationModels.loc[i]["Model"].model.exog_names
```

C:\Users\au479931\AppData\Local\Continuum\anaconda3\lib\site-packages\ipykernel\_launcher.py:29

bla bla

```
In [11]: crossValidationMeanErrors = crossValidationErrors.mean(axis = 1)
        plt.figure()
```

```
plt.plot(crossValidationMeanErrors)
plt.xlabel('Predictors')
plt.ylabel('RSS')
plt.show()
```



bla bla

```
In [12]: print(crossValidationModels.loc[5, "Model"].summary())
```

```

                        OLS Regression Results
=====
Dep. Variable:          Salary    R-squared:                0.771
Model:                  OLS      Adj. R-squared:             0.766
Method:                 Least Squares    F-statistic:          155.0
Date:                  Fri, 04 May 2018    Prob (F-statistic):    1.40e-71
Time:                  17:52:26    Log-Likelihood:        -1704.1
No. Observations:      235    AIC:                   3418.
Df Residuals:          230    BIC:                   3435.
Df Model:               5
Covariance Type:       nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
Walks	6.3470	0.934	6.795	0.000	4.507	8.187
CHits	1.3658	0.353	3.869	0.000	0.670	2.061
CAtBat	-0.2803	0.103	-2.713	0.007	-0.484	-0.077



PutOuts	0.2629	0.083	3.167	0.002	0.099	0.426
CWalks	-0.1555	0.228	-0.681	0.496	-0.605	0.294
=====						
Omnibus:		95.002	Durbin-Watson:			2.014
Prob(Omnibus):		0.000	Jarque-Bera (JB):			564.877
Skew:		1.469	Prob(JB):			2.18e-123
Kurtosis:		10.004	Cond. No.			153.
=====						

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.