

6.6.2 The Lasso

May 18, 2018

```
In [1]: # conventional way to import pandas
import pandas as pd
# conventional way to import seaborn
import seaborn as sns
# conventional way to import numpy
import numpy as np

from sklearn import metrics
import matplotlib.pyplot as plt

data = pd.read_csv("https://vincentarelbundock.github.io/Rdatasets/csv/ISLR/Hitters.csv")

data.head()
```

```
Out[1]:
```

	AtBat	Hits	HmRun	Runs	RBI	Walks	Years	CAtBat	CHits	\
-Andy Allanson	293	66	1	30	29	14	1	293	66	
-Alan Ashby	315	81	7	24	38	39	14	3449	835	
-Alvin Davis	479	130	18	66	72	76	3	1624	457	
-Andre Dawson	496	141	20	65	78	37	11	5628	1575	
-Andres Galarraga	321	87	10	39	42	30	2	396	101	

	CHmRun	CRuns	CRBI	CWalks	League	Division	PutOuts	\
-Andy Allanson	1	30	29	14	A	E	446	
-Alan Ashby	69	321	414	375	N	W	632	
-Alvin Davis	63	224	266	263	A	W	880	
-Andre Dawson	225	828	838	354	N	E	200	
-Andres Galarraga	12	48	46	33	N	E	805	

	Assists	Errors	Salary	NewLeague
-Andy Allanson	33	20	NaN	A
-Alan Ashby	43	10	475.0	N
-Alvin Davis	82	14	480.0	A
-Andre Dawson	11	3	500.0	N
-Andres Galarraga	40	4	91.5	N

After listing the data we can see that some have missing data for their Salary. Next drop all the rows that contain NaN data.

```
In [2]: data = data.dropna()
data.index.name = 'Player'
data.head()
```

```
Out[2]:
```

	AtBat	Hits	HmRun	Runs	RBI	Walks	Years	CAtBat	CHits	\
Player										
-Alan Ashby	315	81	7	24	38	39	14	3449	835	
-Alvin Davis	479	130	18	66	72	76	3	1624	457	
-Andre Dawson	496	141	20	65	78	37	11	5628	1575	
-Andres Galarraga	321	87	10	39	42	30	2	396	101	
-Alfredo Griffin	594	169	4	74	51	35	11	4408	1133	

	CHmRun	CRuns	CRBI	CWalks	League	Division	PutOuts	\
Player								
-Alan Ashby	69	321	414	375	N	W	632	
-Alvin Davis	63	224	266	263	A	W	880	
-Andre Dawson	225	828	838	354	N	E	200	
-Andres Galarraga	12	48	46	33	N	E	805	
-Alfredo Griffin	19	501	336	194	A	W	282	

	Assists	Errors	Salary	NewLeague
Player				
-Alan Ashby	43	10	475.0	N
-Alvin Davis	82	14	480.0	A
-Andre Dawson	11	3	500.0	N
-Andres Galarraga	40	4	91.5	N
-Alfredo Griffin	421	25	750.0	A

In the lab they use a R funktion that we don't have in python. That funktion automatically transforms any qualitative variables into dummy variables. But in python we will have to do this by hand using the following code and display infomation about the variables we converted from strings to numbers.

```
In [3]: dummieVariables = pd.get_dummies(data[['League', 'Division', 'NewLeague']])
dummieVariables.info()
print(dummieVariables.head())
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 263 entries, -Alan Ashby to -Willie Wilson
Data columns (total 6 columns):
League_A      263 non-null uint8
League_N      263 non-null uint8
Division_E    263 non-null uint8
Division_W    263 non-null uint8
NewLeague_A   263 non-null uint8
NewLeague_N   263 non-null uint8
dtypes: uint8(6)
memory usage: 3.6+ KB
```

League_A	League_N	Division_E	Division_W	NewLeague_A	\
----------	----------	------------	------------	-------------	---

Player					
-Alan Ashby	0	1	0	1	0
-Alvin Davis	1	0	0	1	1
-Andre Dawson	0	1	1	0	0
-Andres Galarraga	0	1	1	0	0
-Alfredo Griffin	1	0	0	1	1

	NewLeague_N
Player	
-Alan Ashby	1
-Alvin Davis	0
-Andre Dawson	1
-Andres Galarraga	1
-Alfredo Griffin	0

Next we must remove the columns with our independent variable (Salary), and columns for which we created dummy variables and reintroduce them into our predictors. This is do so the data fit the data in the book.

```
In [4]: y = data.Salary
X_ = data.drop(['Salary', 'League', 'Division', 'NewLeague'], axis=1).astype('float64')

X = pd.concat([X_, dummieVariables[['League_N', 'Division_W', 'NewLeague_N']]], axis=1)
X.head()
```

```
Out[4]:
```

	AtBat	Hits	HmRun	Runs	RBI	Walks	Years	CAtBat	\
Player									
-Alan Ashby	315.0	81.0	7.0	24.0	38.0	39.0	14.0	3449.0	
-Alvin Davis	479.0	130.0	18.0	66.0	72.0	76.0	3.0	1624.0	
-Andre Dawson	496.0	141.0	20.0	65.0	78.0	37.0	11.0	5628.0	
-Andres Galarraga	321.0	87.0	10.0	39.0	42.0	30.0	2.0	396.0	
-Alfredo Griffin	594.0	169.0	4.0	74.0	51.0	35.0	11.0	4408.0	

	CHits	CHmRun	CRuns	CRBI	CWalks	PutOuts	Assists	\
Player								
-Alan Ashby	835.0	69.0	321.0	414.0	375.0	632.0	43.0	
-Alvin Davis	457.0	63.0	224.0	266.0	263.0	880.0	82.0	
-Andre Dawson	1575.0	225.0	828.0	838.0	354.0	200.0	11.0	
-Andres Galarraga	101.0	12.0	48.0	46.0	33.0	805.0	40.0	
-Alfredo Griffin	1133.0	19.0	501.0	336.0	194.0	282.0	421.0	

	Errors	League_N	Division_W	NewLeague_N
Player				
-Alan Ashby	10.0	1	1	1
-Alvin Davis	14.0	0	1	0
-Andre Dawson	3.0	1	0	1
-Andres Galarraga	4.0	1	0	1
-Alfredo Griffin	25.0	0	1	0

Now we split the data into a training and testing set like the book.

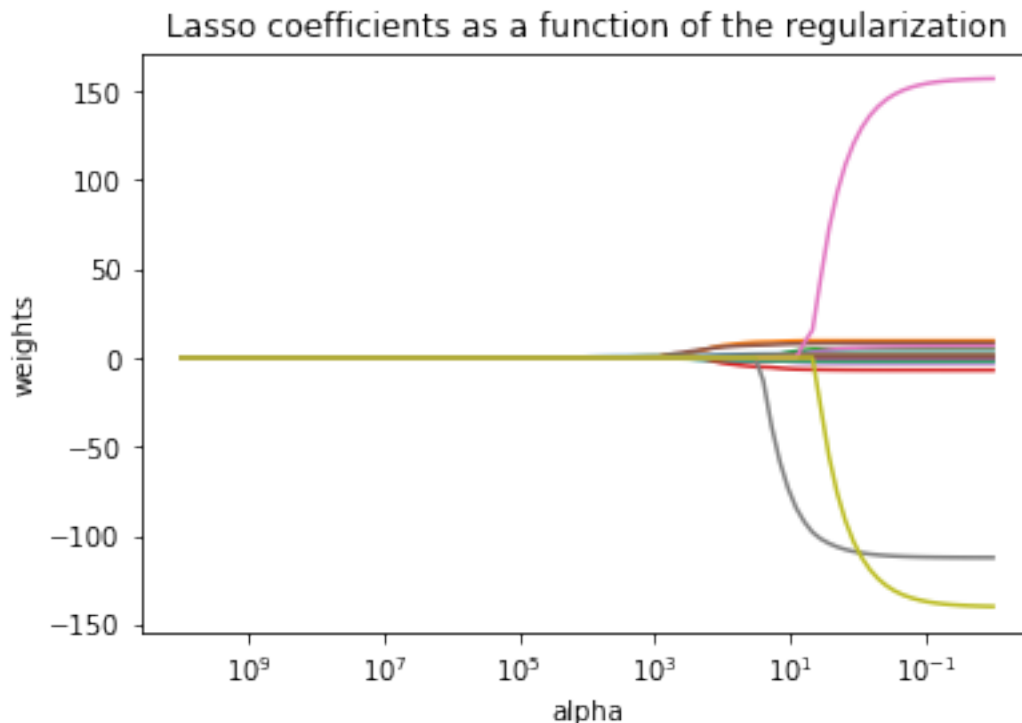
```
In [5]: from sklearn.model_selection import train_test_split
        X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=10)

In [6]: ##Plotting code from scikit-learn.org just changed to use lasso.
        ## http://scikit-learn.org/stable/auto_examples/linear_model/plot_lasso_path.html#sphx
        from sklearn.preprocessing import scale
        from sklearn.linear_model import Lasso

        alphas = 10**np.linspace(10,-2,100)*0.5
        lasso = Lasso(max_iter=100000)
        coefs = []

        for a in alphas*2:
            lasso.set_params(alpha=a)
            lasso.fit(X_train, y_train)
            coefs.append(lasso.coef_)

        ax = plt.gca()
        ax.plot(alphas*2, coefs)
        ax.set_xscale('log')
        ax.set_xlim(ax.get_xlim()[::-1]) # reverse axis
        plt.axis('tight')
        plt.xlabel('alpha')
        plt.ylabel('weights')
        plt.title('Lasso coefficients as a function of the regularization');
```



Next we try to get the best lamda using cross validation.

```
In [7]: from sklearn.linear_model import LassoCV
        alphas = 10**np.linspace(100,-20,100)
        lassoCV = LassoCV(alphas=alphas)
        lassoCV.fit(X_train, y_train)
        lassoCV.alpha_
```

```
C:\Users\entvex\Anaconda3\lib\site-packages\sklearn\linear_model\coordinate_descent.py:491: ConvergenceWarning
```

```
Out [7]: 1072.2672220103254
```

```
In [8]: pd.Series(lassoCV.coef_.flatten(), index=X.columns)
```

```
Out [8]: AtBat          0.729684
        Hits           0.000000
        HmRun          -0.000000
        Runs           0.000000
        RBI            0.000000
        Walks          0.000000
        Years          -0.000000
        CAtBat         -0.241005
        CHits          0.705141
        CHmRun         0.000000
        CRuns          0.579498
        CRBI           0.442905
        CWalks         -0.000000
        PutOuts        0.232895
        Assists        0.000000
        Errors         -0.000000
        League_N       0.000000
        Division_W     -0.000000
        NewLeague_N    0.000000
        dtype: float64
```