

Multiple Linear Regression

May 18, 2018

```
In [5]: # conventional way to import pandas
import pandas as pd
# conventional way to import seaborn
import seaborn as sns
# conventional way to import numpy
import numpy as np

from sklearn import metrics
import matplotlib.pyplot as plt

data = pd.read_csv("https://raw.githubusercontent.com/vincentarelbundock/Rdatasets/master/csv/MASS/


# create a Python list of feature names
feature_cols = ['lstat', 'age']

# use the list to select a subset of the original DataFrame
X = data[feature_cols].values

# select a Series from the DataFrame
y = data['medv'].values
```

Viewing the data in seaborn to get a sense of it - MEDV Median value of owner-occupied homes in \$1000's - LSTAT lower status of the population - AGE proportion of owner-occupied units built prior to 1940

```
In [6]: # visualize the relationship between the features and the response using scatterplots
#sns.pairplot(data, x_vars=['lstat', 'age'], y_vars='medv', size=7, aspect=0.7, kind='r
sns.pairplot(data, x_vars=['lstat', 'age'], y_vars='medv', size=7, aspect=0.7, kind='re
plt.show()
```



```
In [7]: # import model
        from sklearn.linear_model import LinearRegression

        # instantiate
        linreg = LinearRegression()

        # fit the model to the training data (learn the coefficients)
        linreg.fit(X, y)

        # print the intercept and coefficients
        print('intercept: \n', linreg.intercept_)

        # The coefficients
        print('Coefficients: \n', # pair the feature names with the coefficients
              list(zip(feature_cols, linreg.coef_)))

        # make predictions on the testing set
        y_pred = linreg.predict(X)

        # calculate RMSE using scikit-learn
        np.sqrt(metrics.mean_squared_error(y, y_pred))
```

```
# Explained variance score: 1 is perfect prediction
print('Variance score: %.2f' % metrics.r2_score(y, y_pred))
```

intercept:

33.2227605318

Coefficients:

[('lstat', -1.0320685641826013), ('age', 0.034544338571646085)]

Variance score: 0.55

The obtained the predictions produced by linear regression. Look at the residuals, the difference between the real target set and the predicted target set:

```
In [8]: pd.Series(y - y_pred).hist(bins=50)
plt.show()
```

