

Simple linear regression

May 18, 2018


```
In [53]: # conventional way to import pandas
import pandas as pd
# conventional way to import seaborn
import seaborn as sns
# conventional way to import numpy
import numpy as np

from sklearn import metrics
import matplotlib.pyplot as plt

data = pd.read_csv("https://raw.githubusercontent.com/vincentarelbundock/Rdatasets/master/csv/MASS")
```

Viewing the data in seaborn to get a sense of it - MEDV Median value of owner-occupied homes in \$1000's - LSTAT lower status of the population

```
In [54]: # visualize the relationship between the features and the response using scatterplots
#sns.pairplot(data, x_vars='lstat', y_vars='medv', size=7, aspect=0.7, kind='reg')
sns.pairplot(data[['lstat', 'medv']], size=7)
plt.show()
```



Displaying head to see the data

In [55]: `data.head()`

```
Out[55]:
```

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	\
1	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	
2	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	
3	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	
4	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	
5	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	

	black	lstat	medv
1	396.90	4.98	24.0
2	396.90	9.14	21.6

```
3  392.83   4.03  34.7
4  394.63   2.94  33.4
5  396.90   5.33  36.2
```

```
In [56]: # create a Python list of feature names
feature_cols = ['lstat']

# use the list to select a subset of the original DataFrame
X = data[feature_cols].values

# select a Series from the DataFrame
y = data['medv'].values

# import model
from sklearn.linear_model import LinearRegression

# instantiate
linreg = LinearRegression()

# fit the model to the training data (learn the coefficients)
linreg.fit(X, y)

# print the intercept and coefficients
print('intercept: \n', linreg.intercept_)

# The coefficients
print('Coefficients: \n', linreg.coef_)

# make predictions on the testing set
y_pred = linreg.predict(X)

# calculate RMSE using scikit-learn
np.sqrt(metrics.mean_squared_error(y, y_pred))


# Explained variance score: 1 is perfect prediction
print('Variance score: %.2f' % metrics.r2_score(y, y_pred))
```

```
intercept:
34.5538408794
Coefficients:
[-0.95004935]
Variance score: 0.54
```

```
In [57]: # Plot outputs
plt.scatter(X, y)
plt.plot(X, linreg.predict(X), color='blue', linewidth=3)
```

```
plt.xticks(())
plt.yticks(())

plt.show()
```



Display detailed info about the fit. The book also says that the "there is some evidence of non-linearity." So we should reflect that by refing to the R2 score that should be 0.95 to ge a good fit right.

```
In [58]: import statsmodels.api as sm
         from scipy import stats

         X2 = sm.add_constant(X)
         est = sm.OLS(y, X2)
         est2 = est.fit()
         print(est2.summary())
```

OLS Regression Results

```
=====
Dep. Variable:          y      R-squared:          0.544
Model:                  OLS    Adj. R-squared:       0.543
Method:                 Least Squares    F-statistic:       601.6
Date:                  Sun, 18 Feb 2018    Prob (F-statistic): 5.08e-88
Time:                  20:06:53    Log-Likelihood:    -1641.5
No. Observations:      506    AIC:              3287.
Df Residuals:          504    BIC:              3295.
Df Model:               1
Covariance Type:       nonrobust
```


	coef	std err	t	P> t	[0.025	0.975]
const	34.5538	0.563	61.415	0.000	33.448	35.659
x1	-0.9500	0.039	-24.528	0.000	-1.026	-0.874
Omnibus:		137.043	Durbin-Watson:			0.892
Prob(Omnibus):		0.000	Jarque-Bera (JB):			291.373
Skew:		1.453	Prob(JB):			5.36e-64
Kurtosis:		5.319	Cond. No.			29.7

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

The obtained the predictions produced by linear regression. Look at the residuals, the difference between the real target set and the predicted target set:

```
In [59]: pd.Series(y - y_pred).hist(bins=50)
plt.show()
```



Q-Q plot. As it shows in the plot the data is not normal distributed.

```
In [61]: from scipy.stats import probplot
f = plt.figure(figsize=(7, 5))
ax = f.add_subplot(111)
```

```
tuple_out = probplot(y - y_pred, plot=ax)  
plt.show()
```

